

PRESENTATION OUTLINE: A parallel implementation of outlier detection over data streams

Kangqing YU
School of Computer Science
Carleton University
Ottawa, Canada K1S 5B6
kangqingyu@cmail.carleton.ca

November 20, 2017

1 Outlier detection task

- Outlier detection definition
- Typical methods & algorithm
- Main application

2 Data streams

- Characteristics of data streams
- Requirements of real-time data streams processing
- Outlier detection over data streams
- Main challenges

3 Introduction to parallel computing

- GPU and SIMD
- NVIDIA & CUDA framework
- Others

4 Goal of this project

- Intersection of outlier detection, data streams and parallel computations
- Brief description of proposed approach
- Novelty

5 Local Outlier Factor(LOF) algorithm

- General idea
- Steps and mathematical formulas
- Extremely computational expensive

6 Statistical binned summary

- What is contained in the binned summary
- How to calculate the bin index for each data point
- How to maintain/update the binned summary
- Curse of dimension problem

7 Cumulative LOF algorithm(modified version)

- General idea
- Role of the binned summary
- Steps and mathematical formulas

8 Implementation of cumulative LOF on CUDA

- Coalesced memory access & shared memory
- Detailed steps in host and kernel functions
- Threads managements in each kernels

9 Binned summary maintenance using MapReduce

- Why MapReduce?
- Map stage
- Reduce stage

10 Q & A

- Questions?